

EDA on used cars

November 27, 2022

```
[1]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: # Getting the current working directory
os.getcwd()
```

```
[2]: 'C:\\Users\\shoaib\\Desktop\\Revise Py-2022'
```

```
[3]: #Reading the data from csv
cars_data=pd.read_csv("cars_sampled.csv")
```

```
[4]: #Creating copy of cars_data
cars=cars_data.copy() # deep copy
#Whatever changes made in cars will not be reflected in cars_data
```

```
[5]: cars.head()
```

```
[5]:      dateCrawled      name \
0  30/03/2016 13:51      Zu_verkaufen
1    7/3/2016 9:54  Volvo_XC90_2.4D_Summum
2   1/4/2016 0:57  Volkswagen_Touran
3  19/03/2016 17:50  Seat_Ibiza_1.4_16V_Reference
4  16/03/2016 14:51  Volvo_XC90_D5_Aut._RDesign_R_Design_AWD_GSHD_S...

      seller offerType  price  abtest vehicleType  yearOfRegistration \
0  private      offer   4450   test   limousine              2003
1  private      offer  13299  control         suv              2005
2  private      offer   3200   test         bus              2003
3  private      offer   4500  control   small car              2006
4  private      offer  18750   test         suv              2008

      gearbox  powerPS      model  kilometer  monthOfRegistration  fuelType \
0    manual     150        3er    150000              3    diesel
1    manual     163  xc_reihe    150000              6    diesel
2    manual     101    touran    150000             11    diesel
3    manual      86    ibiza     60000             12    petrol
4  automatic     185  xc_reihe    150000             11    diesel
```

| | brand | notRepairedDamage | dateCreated | postalCode | lastSeen |
|---|------------|-------------------|-----------------|------------|------------------|
| 0 | bmw | NaN | 30/03/2016 0:00 | 20257 | 7/4/2016 4:44 |
| 1 | volvo | no | 7/3/2016 0:00 | 88045 | 26/03/2016 13:17 |
| 2 | volkswagen | NaN | 31/03/2016 0:00 | 27449 | 1/4/2016 8:40 |
| 3 | seat | no | 19/03/2016 0:00 | 34537 | 7/4/2016 4:44 |
| 4 | volvo | no | 16/03/2016 0:00 | 55270 | 1/4/2016 23:18 |

```
[6]: cars.info()
#This gives the total entries in a column and its data type
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50001 entries, 0 to 50000
Data columns (total 19 columns):
dateCrawled      50001 non-null object
name             50001 non-null object
seller           50001 non-null object
offerType        50001 non-null object
price            50001 non-null int64
abtest           50001 non-null object
vehicleType      44813 non-null object
yearOfRegistration 50001 non-null int64
gearbox          47177 non-null object
powerPS          50001 non-null int64
model            47243 non-null object
kilometer        50001 non-null int64
monthOfRegistration 50001 non-null int64
fuelType         45498 non-null object
brand            50001 non-null object
notRepairedDamage 40285 non-null object
dateCreated      50001 non-null object
postalCode       50001 non-null int64
lastSeen         50001 non-null object
dtypes: int64(6), object(13)
memory usage: 7.2+ MB
```

```
[7]: pd.set_option('display.float_format', lambda x: '%.3f'%x) # Get the floating_
      ↪point to 3 decimal places
pd.set_option('display.max_columns', 500) # To display maximum set of columns
cars.describe() #This gives the statistical info about the data- https://www.
      ↪mathsisfun.com/data/standard-deviation.html
# https://statisticsbyjim.com/basics/coefficient-variation/
```

```
[7]:
```

| | price | yearOfRegistration | powerPS | kilometer | \ |
|-------|-----------|--------------------|-----------|------------|---|
| count | 50001.000 | 50001.000 | 50001.000 | 50001.000 | |
| mean | 6559.865 | 2005.544 | 116.496 | 125613.688 | |
| std | 85818.470 | 122.992 | 230.568 | 40205.234 | |
| min | 0.000 | 1000.000 | 0.000 | 5000.000 | |

| | | | | |
|-----|--------------|----------|-----------|------------|
| 25% | 1150.000 | 1999.000 | 69.000 | 125000.000 |
| 50% | 2950.000 | 2003.000 | 105.000 | 150000.000 |
| 75% | 7190.000 | 2008.000 | 150.000 | 150000.000 |
| max | 12345678.000 | 9999.000 | 19312.000 | 150000.000 |

| | | |
|-------|---------------------|------------|
| | monthOfRegistration | postalCode |
| count | 50001.000 | 50001.000 |
| mean | 5.744 | 50775.217 |
| std | 3.711 | 25743.702 |
| min | 0.000 | 1067.000 |
| 25% | 3.000 | 30559.000 |
| 50% | 6.000 | 49504.000 |
| 75% | 9.000 | 71404.000 |
| max | 12.000 | 99998.000 |

```
[8]: print('-----')
      print('          Data Cleaning          ')
      print('-----')
```

```
-----
          Data Cleaning
-----
```

```
[9]: #Dropping unwanted columns as these variables are related ads
      cars=cars.drop(['dateCrawled','name','dateCreated','lastSeen'],axis=1)
```

```
[10]: #Removing the duplicates
       cars.drop_duplicates(keep='first',inplace=True)
       cars.info()
       #16 duplicates removed from the cars dataset
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 49666 entries, 0 to 50000
Data columns (total 15 columns):
seller                49666 non-null object
offerType             49666 non-null object
price                 49666 non-null int64
abtest                49666 non-null object
vehicleType           44491 non-null object
yearOfRegistration    49666 non-null int64
gearbox               46855 non-null object
powerPS               49666 non-null int64
model                 46919 non-null object
kilometer             49666 non-null int64
monthOfRegistration    49666 non-null int64
fuelType              45177 non-null object
brand                 49666 non-null object
notRepairedDamage     39980 non-null object
```

```
postalCode          49666 non-null int64
dtypes: int64(6), object(9)
memory usage: 6.1+ MB
```

```
[11]: #Checking the missing values from each variable
cars.isnull().sum()
```

```
[11]: seller          0
offerType          0
price             0
abtest            0
vehicleType       5175
yearOfRegistration 0
gearbox          2811
powerPS           0
model            2747
kilometer         0
monthOfRegistration 0
fuelType         4489
brand             0
notRepairedDamage 9686
postalCode        0
dtype: int64
```

```
[12]: #variable- year of registration
print(cars['yearOfRegistration'].value_counts().sort_index())
#There are some cars which are registered before 1900 and there are also cars
→after 2019
```

```
1000      6
1255      1
1500      2
1910     15
1928      1
1929      1
1933      1
1934      1
1936      2
1938      1
1940      1
1941      1
1943      2
1945      2
1947      2
1950      4
1951      4
1952      3
1953      2
```

| | |
|------|------|
| 1954 | 1 |
| 1955 | 6 |
| 1956 | 7 |
| 1957 | 5 |
| 1958 | 4 |
| 1959 | 5 |
| 1960 | 33 |
| 1961 | 7 |
| 1962 | 6 |
| 1963 | 11 |
| 1964 | 16 |
| ... | |
| 2002 | 2560 |
| 2003 | 2742 |
| 2004 | 2614 |
| 2005 | 3109 |
| 2006 | 2668 |
| 2007 | 2357 |
| 2008 | 2190 |
| 2009 | 2018 |
| 2010 | 1645 |
| 2011 | 1548 |
| 2012 | 1235 |
| 2013 | 821 |
| 2014 | 624 |
| 2015 | 406 |
| 2016 | 1351 |
| 2017 | 1376 |
| 2018 | 528 |
| 2019 | 2 |
| 2222 | 1 |
| 2900 | 1 |
| 3000 | 1 |
| 3500 | 1 |
| 3800 | 1 |
| 5000 | 3 |
| 6000 | 4 |
| 7500 | 1 |
| 7800 | 1 |
| 8500 | 1 |
| 8888 | 2 |
| 9999 | 7 |


Name: yearOfRegistration, Length: 97, dtype: int64

```
[13]: print(cars['yearOfRegistration'].describe())
      #There is no much difference b/w mean and median but std dev is differing from
      →mean by 123 years
```

```
sns.
→regplot(x='yearOfRegistration',y='price',scatter=True,fit_reg=False,data=cars)
#Here we can't get any clear picture as there lot of outliers. so data is
→concentrated at one point
```

```
count    49666.000
mean      2005.551
std       123.405
min       1000.000
25%      1999.000
50%      2003.000
75%      2008.000
max       9999.000
Name: yearOfRegistration, dtype: float64
```

[13]: <matplotlib.axes._subplots.AxesSubplot at 0x15fda7a3940>



```
[14]: #After trial error, getting working range for year of registration
# Working range- 1950 and 2018
print(sum(cars['yearOfRegistration'] > 2018))
print(sum(cars['yearOfRegistration'] < 1950))
```

26
39

```
[15]: #Variable - Price
cars['price'].value_counts().sort_index()
#There are 1437 cars at 0 dollar
```

```
[15]: 0          1437
      1          172
      2           1
      3           1
      5           4
      7           1
      8           2
     10           5
     11           1
     12           1
     14           1
     15           8
     20           6
     21           1
     25           5
     26           1
     30           7
     35           4
     39           1
     40           3
     45           6
     50          41
     55           3
     60           7
     65           1
     70           2
     75           9
     77           1
     80          12
     85           3

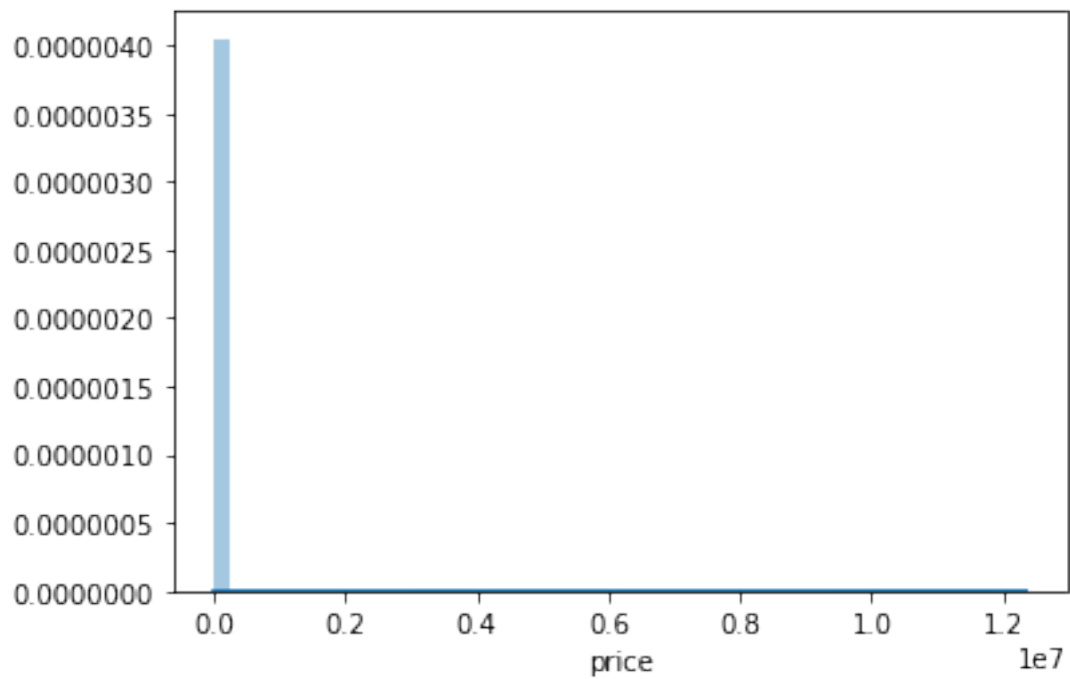
      ...
163991          1
165000          1
169999          1
171000          1
175000          1
179999          1
189981          1
205000          1
214800          1
225000          1
230000          2
239000          1
```

| | |
|----------|---|
| 249000 | 1 |
| 250000 | 1 |
| 257500 | 1 |
| 260000 | 1 |
| 270000 | 1 |
| 300000 | 1 |
| 370000 | 1 |
| 395000 | 1 |
| 485000 | 1 |
| 487000 | 1 |
| 619000 | 1 |
| 700000 | 1 |
| 999999 | 1 |
| 1250000 | 1 |
| 2795000 | 1 |
| 9999999 | 1 |
| 10010011 | 1 |
| 12345678 | 1 |

Name: price, Length: 2393, dtype: int64

```
[16]: print(sns.distplot(cars['price']))
      #Here the values are cluttered at zero
```

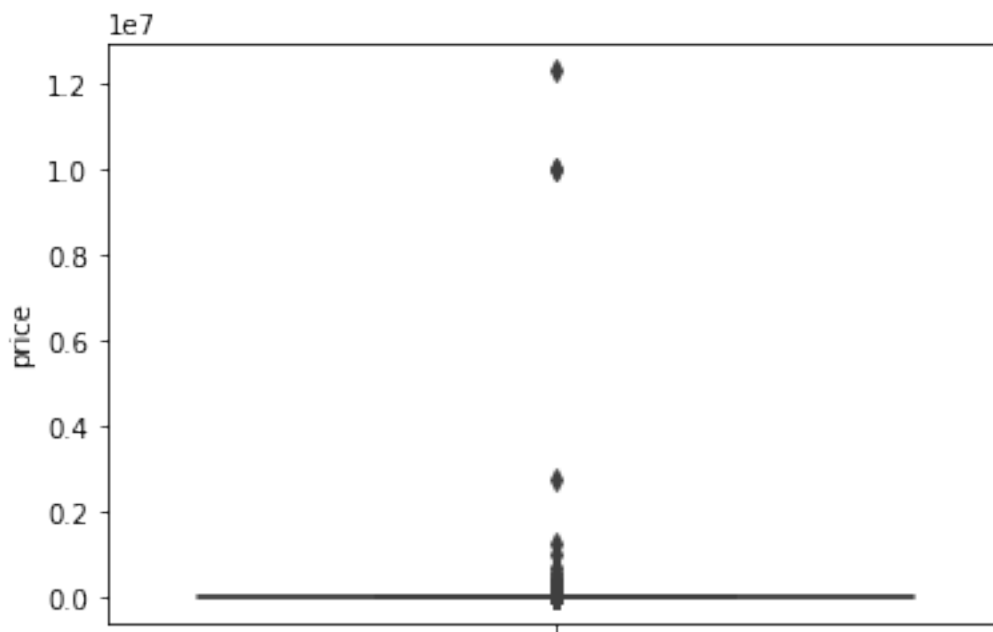
AxesSubplot(0.125,0.125;0.775x0.755)




```
[17]: print(sns.boxplot(y=cars['price']))
      # similarly boxplot is not clearly visible due to outliers
      cars['price'].describe()
      # Std dev is 1300% of mean
```

AxesSubplot(0.125,0.125;0.775x0.755)

```
[17]: count      49666.000
      mean        6559.299
      std         86105.705
      min           0.000
      25%         1150.000
      50%         2950.000
      75%         7100.000
      max        12345678.000
      Name: price, dtype: float64
```



```
[18]: #count
      print(sum(cars['price'] > 150000))
      print(sum(cars['price'] < 100))
      # Working range- 100 and 150000 after trial and error
```

34
1770

```
[19]: # Variable powerPS
cars['powerPS'].value_counts().sort_index()
#5581 cars have zero horse power
```

```
[19]: 0          5581
      1           3
      2           2
      3           2
      4           4
      5          17
      6           2
      7           1
      8           1
      9           2
     10           4
     11           6
     12           1
     13           6
     14           3
     15           1
     16           9
     17           2
     18           1
     19           1
     20           3
     21           5
     22          46
     23           5
     24           2
     25           7
     26           1
     27           2
     28           5
     29           2
     30           29
     ...
    1223           1
    1256           1
    1363           1
    1416           1
    1502           1
    1595           1
    1598           1
    1625           1
    1653           1
    1799           1
    1910           1
    1968           1
```

```

1992      1
1998      1
2004      1
2017      1
2172      1
2461      1
2789      1
6226      1
11620     1
12510     1
12512     1
12684     1
15017     1
15033     1
16011     1
16312     1
19211     1
19312     1
Name: powerPS, Length: 460, dtype: int64

```

```


[20]: sns.regplot(x='powerPS',y='price',scatter=True,fit_reg=False,data=cars)
      #data is concentrated at zero power ps

```

```

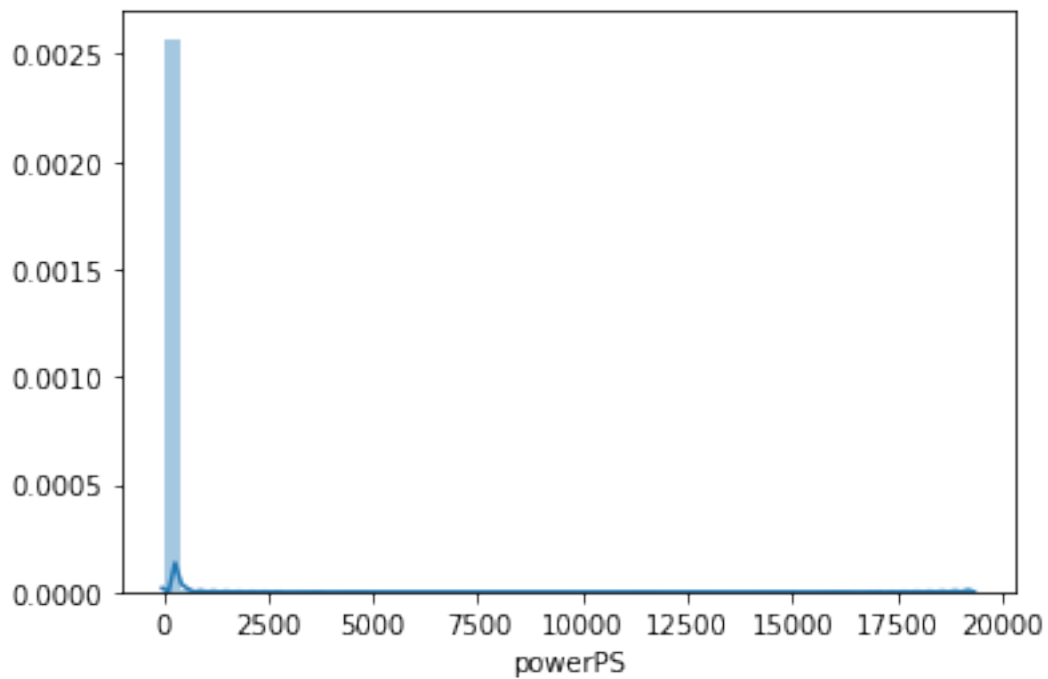
[20]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdbaf9d68>

```



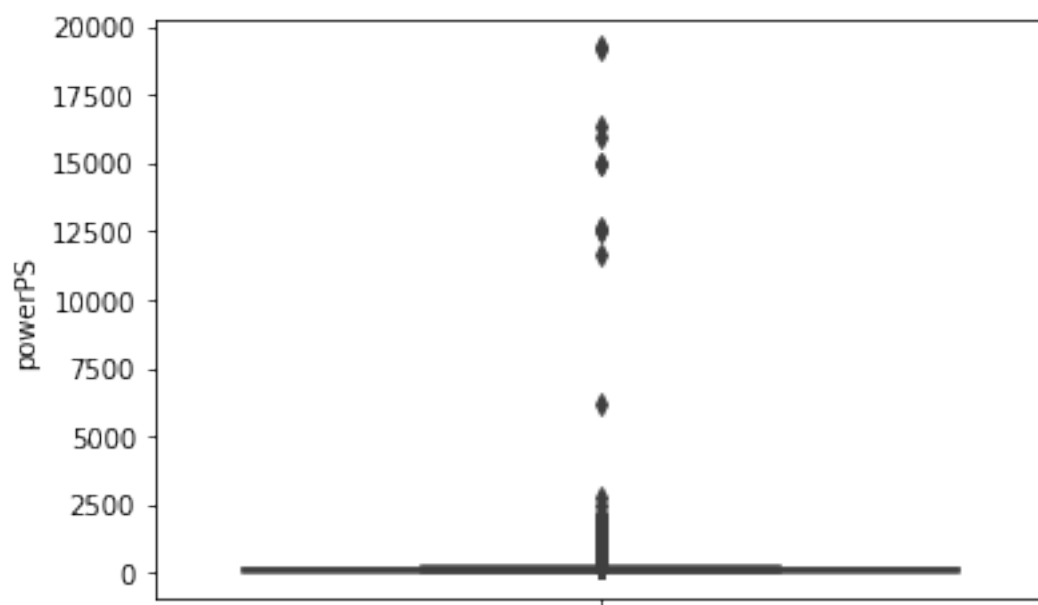
```
[21]: sns.distplot(cars['powerPS'])
```

```
[21]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdbbbe940>
```



```
[22]: sns.boxplot(y=cars['powerPS'])
```

```
[22]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdbcc5f60>
```



```
[23]: cars['powerPS'].describe()
      #std dev is 200% of mean
      #minimum is zero but engine cannot start at that HP
```

```
[23]: count    49666.000
      mean      116.404
      std       231.262
      min        0.000
      25%       69.000
      50%      105.000
      75%      150.000
      max     19312.000
      Name: powerPS, dtype: float64
```

```
[24]: # Working range- 10 and 500
      print(sum(cars['powerPS'] > 500))
      print(sum(cars['powerPS'] < 10))
```

```
115
5613
```

```
[25]: # =====
      # Working range of data
      # =====
      cars=cars[
          (cars.yearOfRegistration <= 2018)
          & (cars.yearOfRegistration >= 1950)
          & (cars.price >= 100)
          & (cars.price <= 150000)
          & (cars.powerPS >= 10)
          & (cars.powerPS <= 500)]
      cars.shape
      #~7300 records are dropped
```

```
[25]: (42855, 15)
```

```
[26]: #Variable reduction
      #Combining year of registration and month of registration to form new variable_
      ↪Age
      cars['monthOfRegistration']/=12
      cars.head()
```

```
[26]:   seller offerType  price  abtest vehicleType  yearOfRegistration  \
0  private    offer   4450    test   limousine           2003
1  private    offer  13299  control         suv           2005
2  private    offer   3200    test         bus           2003
3  private    offer   4500  control   small car           2006
4  private    offer  18750    test         suv           2008
```

| | gearbox | powerPS | model | kilometer | monthOfRegistration | fuelType | \ |
|---|-----------|---------|----------|-----------|---------------------|----------|---|
| 0 | manual | 150 | 3er | 150000 | 0.250 | diesel | |
| 1 | manual | 163 | xc_reihe | 150000 | 0.500 | diesel | |
| 2 | manual | 101 | touran | 150000 | 0.917 | diesel | |
| 3 | manual | 86 | ibiza | 60000 | 1.000 | petrol | |
| 4 | automatic | 185 | xc_reihe | 150000 | 0.917 | diesel | |

| | brand | notRepairedDamage | postalCode |
|---|------------|-------------------|------------|
| 0 | bmw | NaN | 20257 |
| 1 | volvo | no | 88045 |
| 2 | volkswagen | NaN | 27449 |
| 3 | seat | no | 34537 |
| 4 | volvo | no | 55270 |

```
[27]: #new variable Age
cars['Age']=(2018-cars['yearOfRegistration']) + cars['monthOfRegistration']
cars['Age']= round(cars['Age'],2)
```

```
[28]: cars['Age'].describe()
```

```
[28]: count    42855.000
mean         14.872
std           7.090
min           0.000
25%          10.330
50%          14.830
75%          19.170
max           67.750
Name: Age, dtype: float64
```

```
[29]: #dropping year of registration and month of registration
cars=cars.drop(['yearOfRegistration','monthOfRegistration'],axis=1)
```

```
[30]: cars.head()
```

| | seller | offerType | price | abtest | vehicleType | gearbox | powerPS | \ |
|---|---------|-----------|-------|---------|-------------|-----------|---------|---|
| 0 | private | offer | 4450 | test | limousine | manual | 150 | |
| 1 | private | offer | 13299 | control | suv | manual | 163 | |
| 2 | private | offer | 3200 | test | bus | manual | 101 | |
| 3 | private | offer | 4500 | control | small car | manual | 86 | |
| 4 | private | offer | 18750 | test | suv | automatic | 185 | |

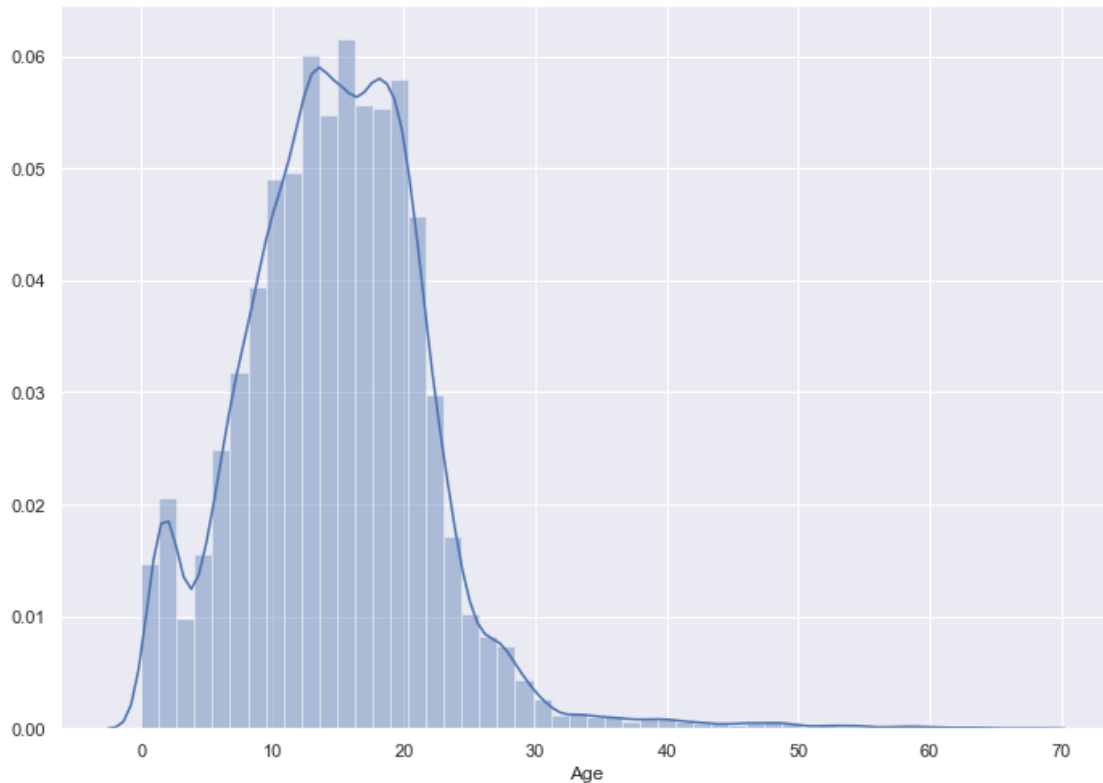
| | model | kilometer | fuelType | brand | notRepairedDamage | postalCode | \ |
|---|----------|-----------|----------|------------|-------------------|------------|---|
| 0 | 3er | 150000 | diesel | bmw | NaN | 20257 | |
| 1 | xc_reihe | 150000 | diesel | volvo | no | 88045 | |
| 2 | touran | 150000 | diesel | volkswagen | NaN | 27449 | |
| 3 | ibiza | 60000 | petrol | seat | no | 34537 | |
| 4 | xc_reihe | 150000 | diesel | volvo | no | 55270 | |

```
Age
0 15.250
1 13.500
2 15.920
3 13.000
4 10.920
```

```
[31]: # =====
# Visualizing Parameters
# =====
sns.set(rc={'figure.figsize':(11.7,8.27)}) #A4 size dimension
```

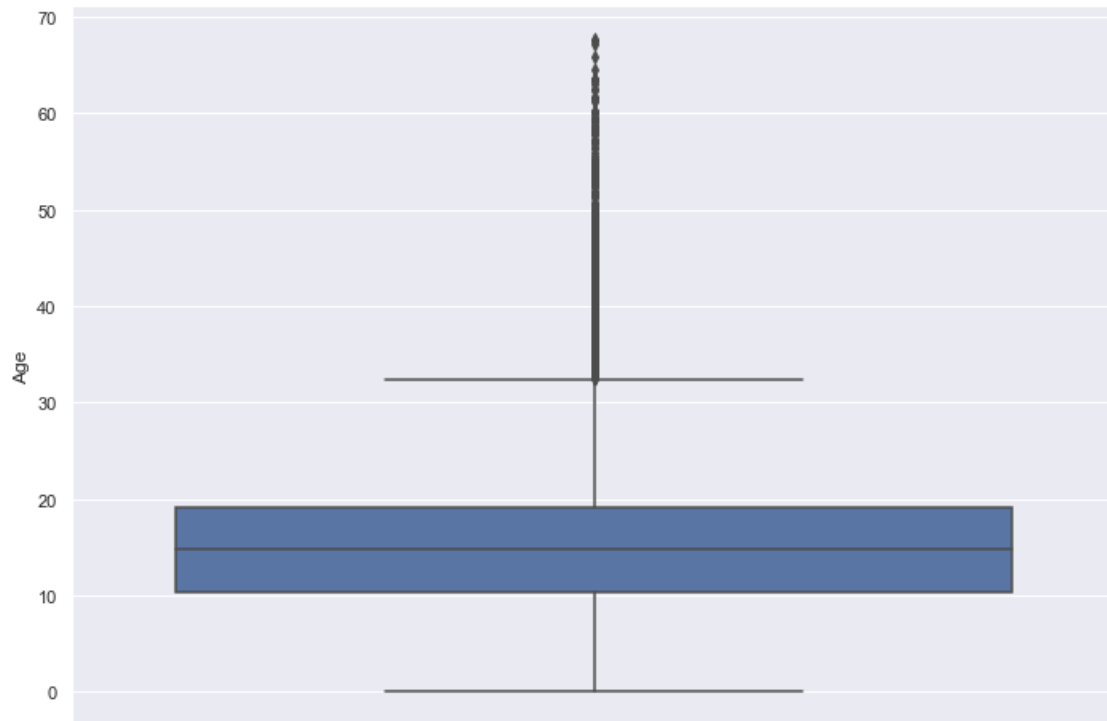
```
[32]: #Age
sns.distplot(cars['Age'])
#data is slightly skewed towards right
```

```
[32]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdbb5d2e8>
```



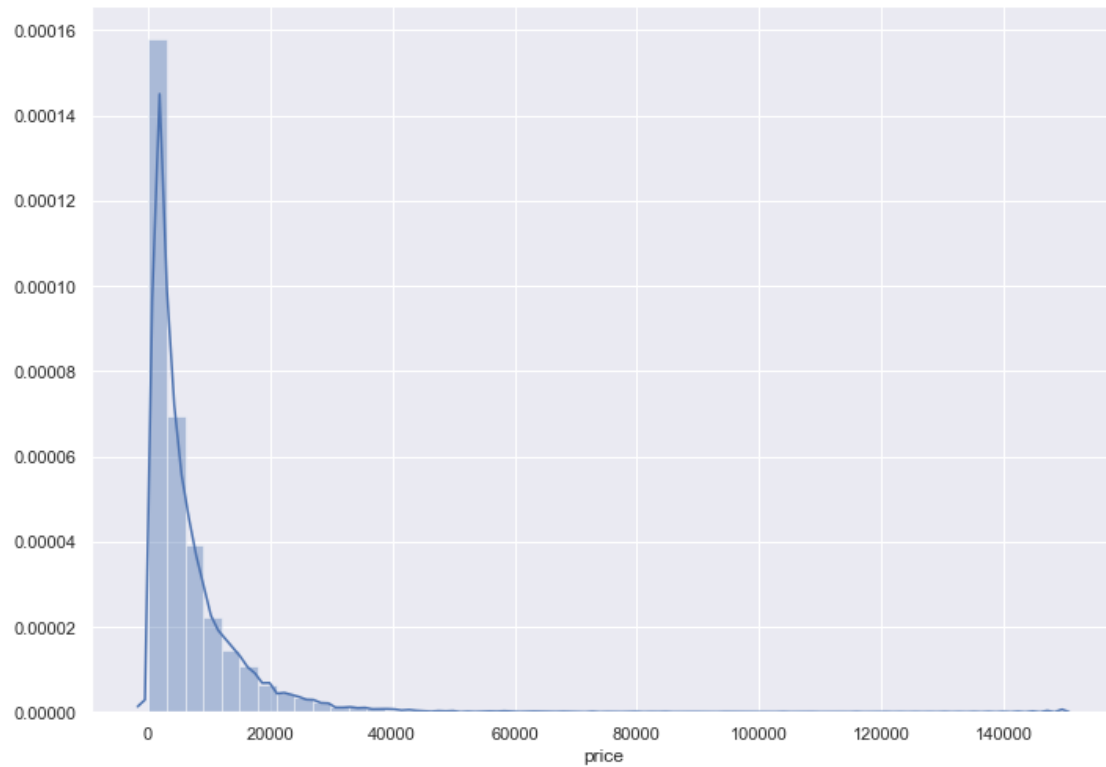
```
[33]: sns.boxplot(y=cars['Age'])
# Here it can be see that plots evenly distributed but still there are few
→outliers
```

```
[33]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdb0860b8>
```



```
[34]: #price
sns.distplot(cars['price'])
#plot is right skewed because of high end luxurious cars
```

```
[34]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdb048cc0>
```

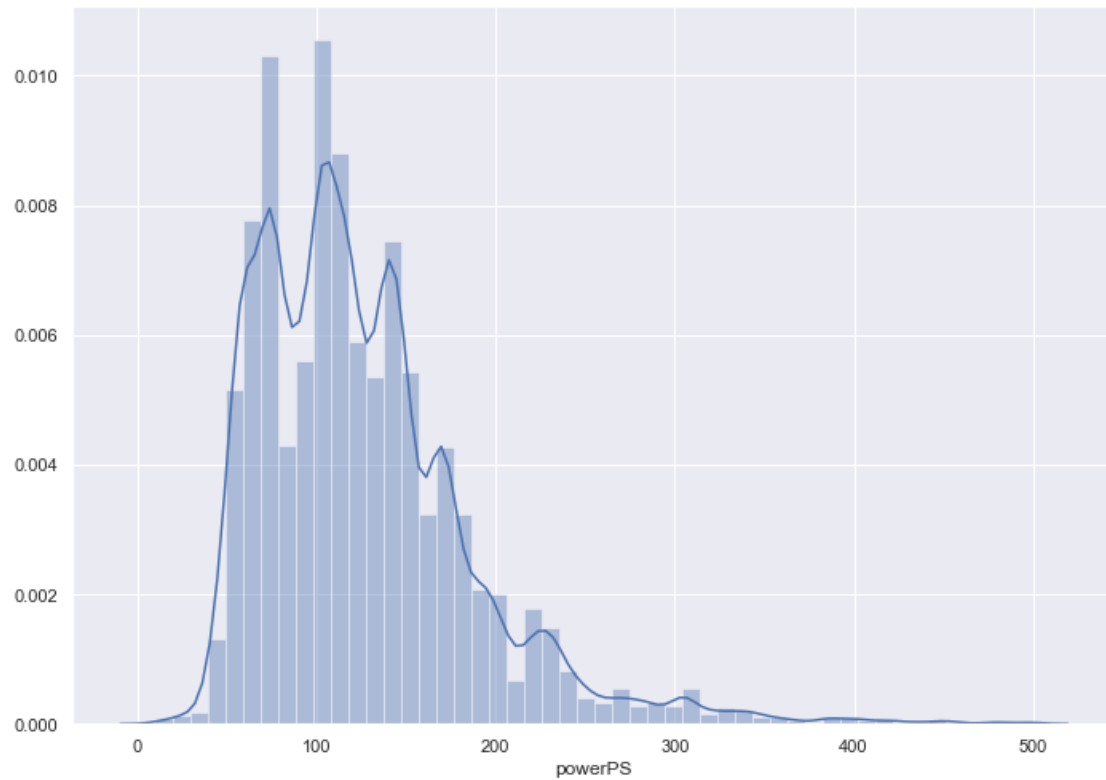
```
[35]: sns.boxplot(y=cars['price'])  
cars['price'].describe()  
#std dev is still high due to luxurious cars
```

```
[35]: count    42855.000  
      mean      6132.950  
      std      7944.135  
      min       100.000  
      25%     1450.000  
      50%     3499.000  
      75%     7900.000  
      max    149000.000  
      Name: price, dtype: float64
```



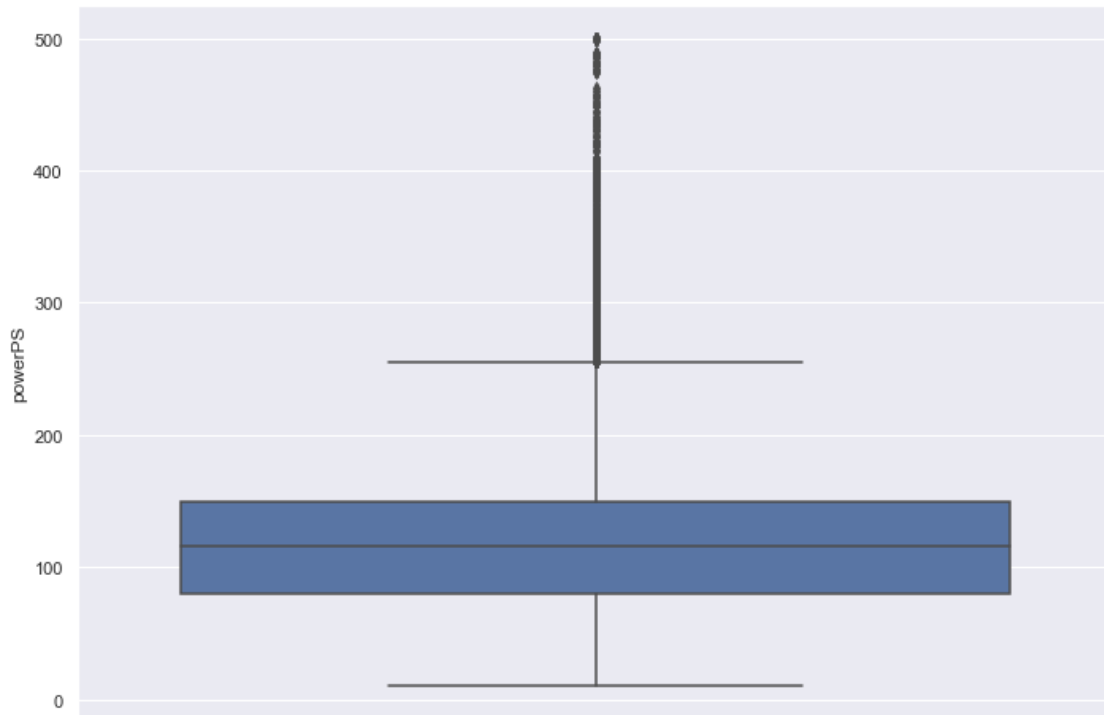
```
[36]: #powerPS  
sns.distplot(cars['powerPS'])
```

```
[36]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdca838d0>
```



```
[37]: sns.boxplot(y=cars['powerPS'])  
cars['powerPS'].describe()
```

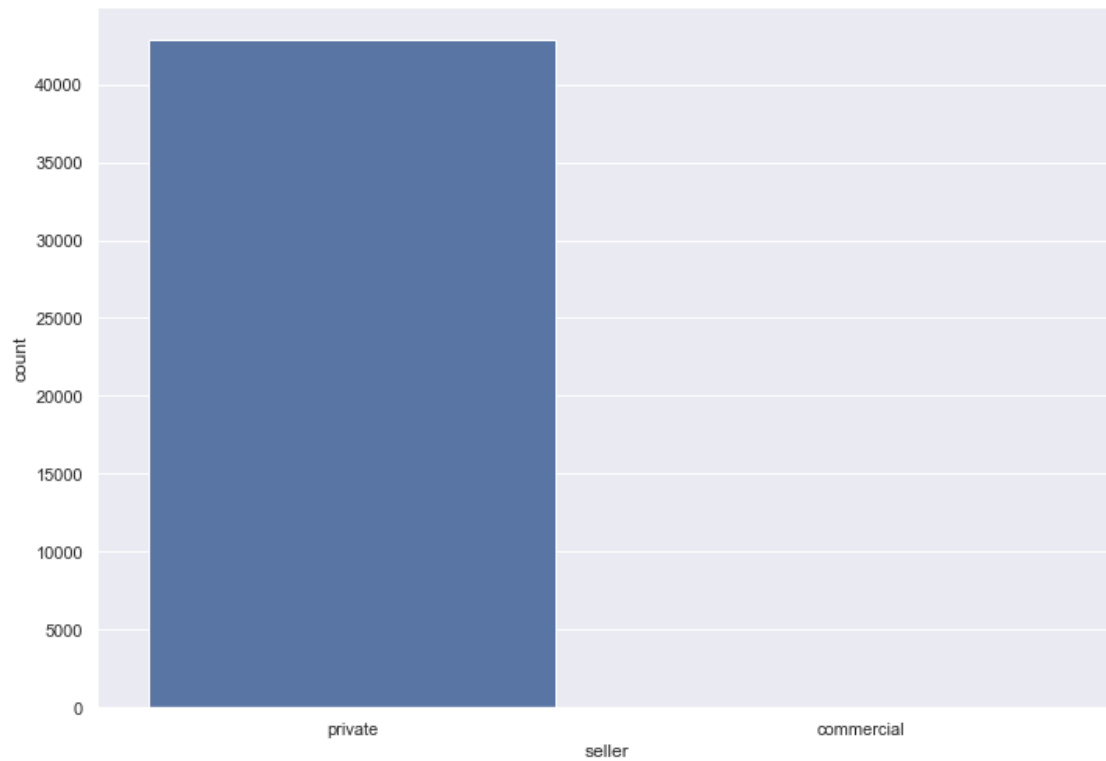
```
[37]: count    42855.000  
      mean      126.054  
      std       60.525  
      min       10.000  
      25%       80.000  
      50%      116.000  
      75%      150.000  
      max      500.000  
      Name: powerPS, dtype: float64
```



```
[38]: #Variable Age
print(cars['seller'].value_counts())
print(sns.countplot(cars['seller']))
pd.crosstab(cars['seller'], columns='count', normalize=True)
#almost all cars from private players; 99.99% of sellers at storm motors are
→private persons not commercial businessess
```

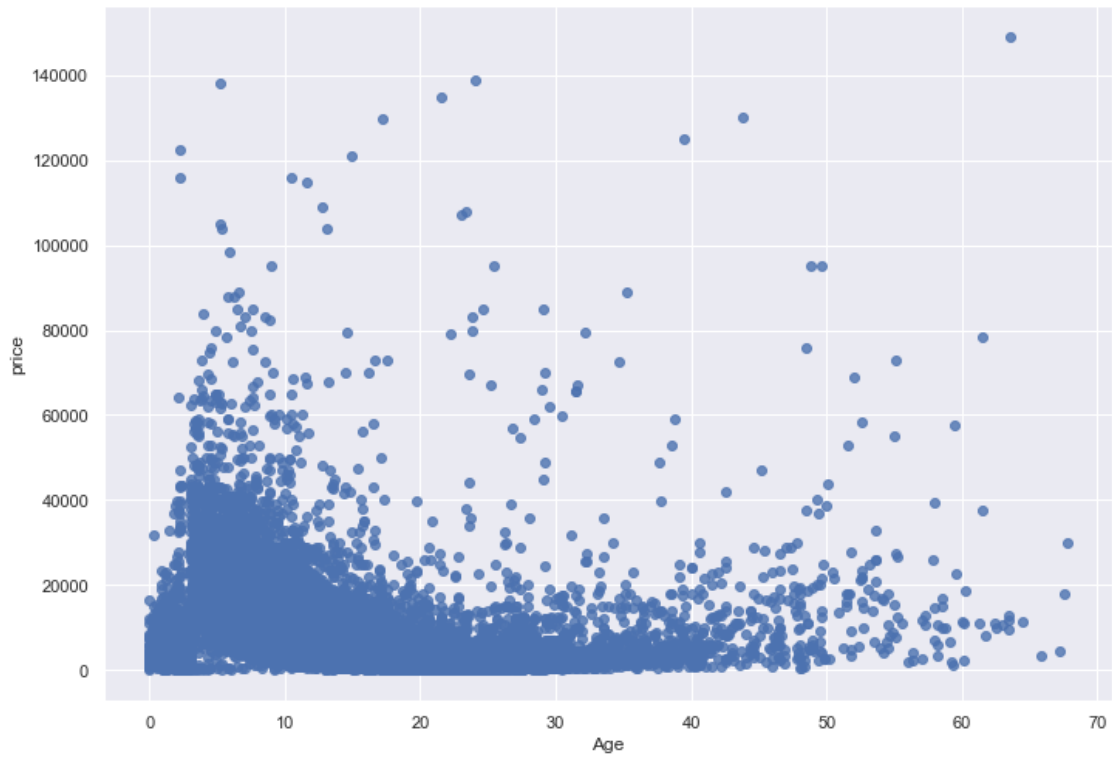
```
private      42854
commercial      1
Name: seller, dtype: int64
AxesSubplot(0.125,0.125;0.775x0.755)
```

```
[38]: col_0      count
seller
commercial  0.000
private    1.000
```



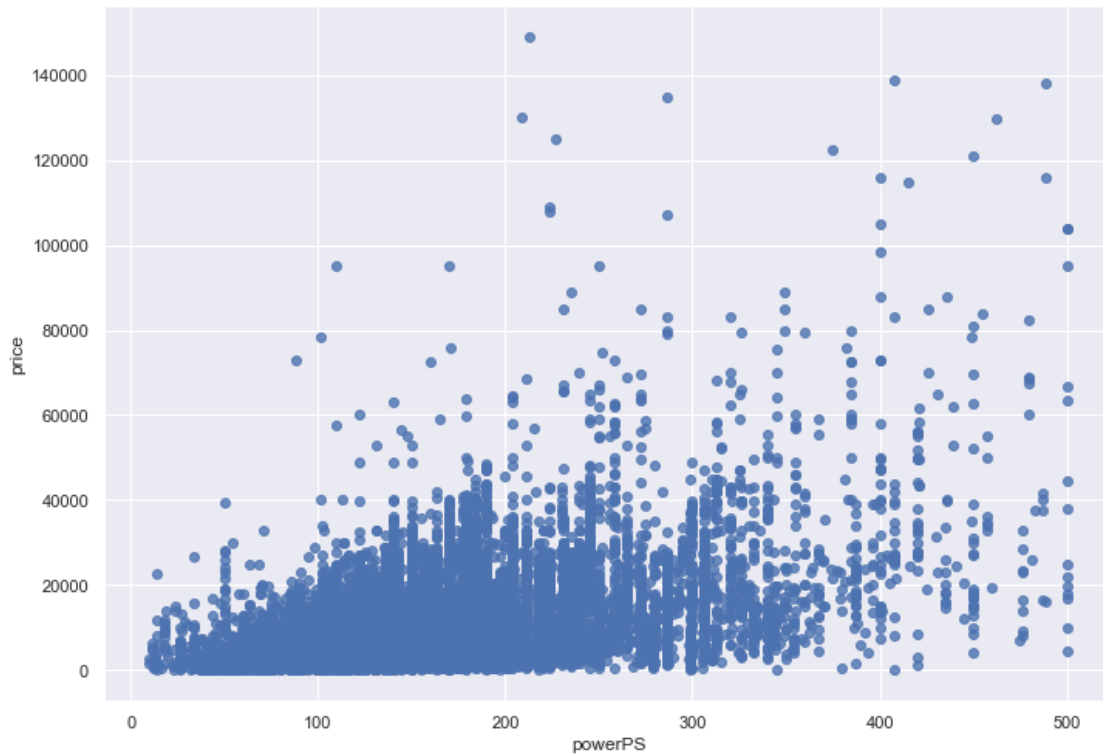
```
[39]: # Age Vs Price
sns.regplot(x='Age',y='price',scatter=True,fit_reg=False,data=cars)
#As the Age of a vehicle increases, price starts to decrease
#but there are some cars whose price is higher despite increase in age
```

```
[39]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdd22c1d0>
```




```
[40]: # powerPS vs price
sns.regplot(x='powerPS', y='price', scatter=True,
            fit_reg=False, data=cars)
#as the power PS increases, price also increases
```

```
[40]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdd296fd0>
```



```
[41]: #Price Vs Kilometer
sns.boxplot(x='kilometer',y='price',data=cars)
#As running of a vehicle increases, price gets reduce except for vehicles below
→5000 km running
#cars within 5000km running having low price because of high age
```

```
[41]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdce94f60>
```



```
[42]: cars_age=cars[(cars['kilometer']<= 5000)]
print(cars_age.shape)
cars_age.describe()
# The Average Age of 479 cars is same as the average age of 42000 cars
#which means cars which have running below 5000km have high aged cars due to
→which price is lower
```

(479, 14)

```
[42]:
```

| | price | powerPS | kilometer | postalCode | Age |
|-------|------------|---------|-----------|------------|---------|
| count | 479.000 | 479.000 | 479.000 | 479.000 | 479.000 |
| mean | 9950.616 | 124.649 | 5000.000 | 49691.161 | 14.844 |
| std | 15939.343 | 72.072 | 0.000 | 25925.783 | 12.320 |
| min | 101.000 | 14.000 | 5000.000 | 1069.000 | 0.170 |
| 25% | 1100.000 | 75.000 | 5000.000 | 28540.500 | 3.750 |
| 50% | 3000.000 | 110.000 | 5000.000 | 47269.000 | 15.170 |
| 75% | 11949.500 | 150.000 | 5000.000 | 70702.000 | 20.500 |
| max | 149000.000 | 500.000 | 5000.000 | 99817.000 | 63.580 |

```
[43]: cars_age1=cars[(cars['kilometer'] > 5000)]
print(cars_age1.shape)
cars_age1.describe()
```

(42376, 14)

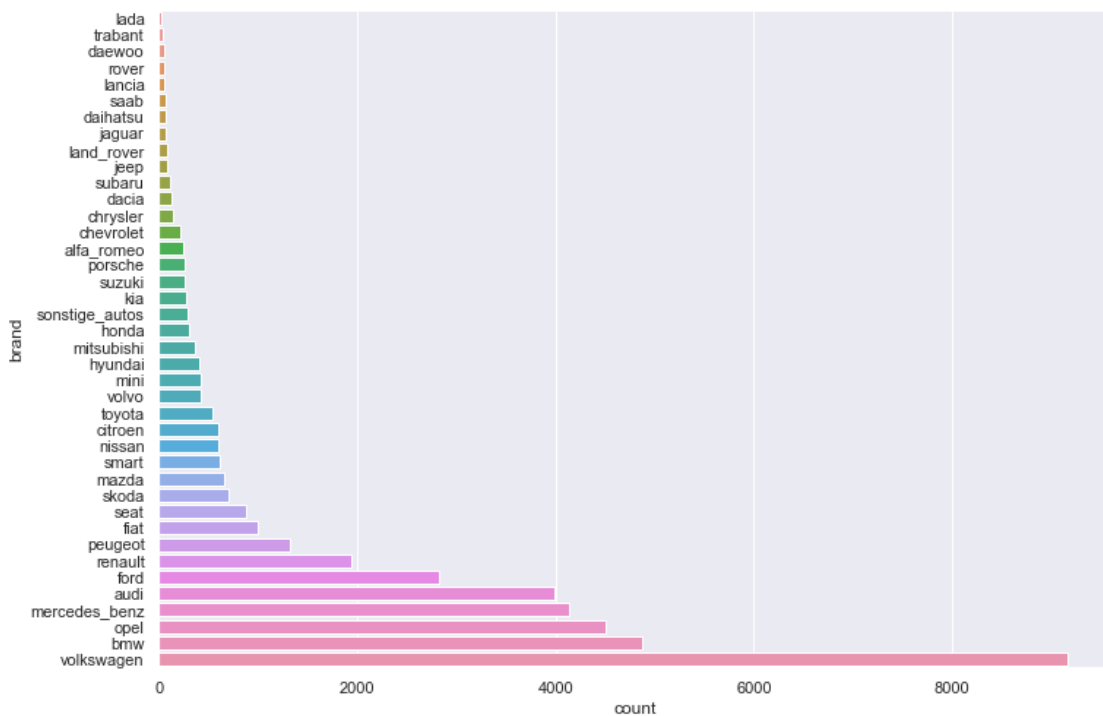

```
[43]:
```

| | price | powerPS | kilometer | postalCode | Age |
|-------|------------|-----------|------------|------------|-----------|
| count | 42376.000 | 42376.000 | 42376.000 | 42376.000 | 42376.000 |
| mean | 6089.797 | 126.070 | 127188.149 | 51576.149 | 14.872 |
| std | 7796.804 | 60.382 | 37106.044 | 25717.238 | 7.008 |
| min | 100.000 | 10.000 | 10000.000 | 1067.000 | 0.000 |
| 25% | 1450.000 | 80.000 | 125000.000 | 31275.000 | 10.330 |
| 50% | 3499.000 | 116.000 | 150000.000 | 50674.000 | 14.830 |
| 75% | 7899.000 | 150.000 | 150000.000 | 72414.000 | 19.170 |
| max | 139000.000 | 500.000 | 150000.000 | 99998.000 | 67.750 |

```
[44]: #cars[(cars['kilometer'] <= 5000) & (cars['Age'] > 10)]
# 293 cars among 500 cars have low price despite running lower than 5000 km but
→age is high due to which price is low
```

```
[45]: sns.countplot(y='brand',data=cars,order=cars['brand'] .
→value_counts(ascending=True).index)
# Volkswagen is highest sold brand and lada is least sold brand to the company
```

```
[45]: <matplotlib.axes._subplots.AxesSubplot at 0x15fdcacce80>
```



```
[46]: pd.crosstab(cars['brand'],columns='count',normalize=True,colnames=['%']).
→sort_values(by='count',ascending=False)*100
```

```
[46]: %          count
brand
volkswagen    21.379
```

| | |
|----------------|--------|
| bmw | 11.383 |
| opel | 10.498 |
| mercedes_benz | 9.665 |
| audi | 9.317 |
| ford | 6.583 |
| renault | 4.539 |
| peugeot | 3.089 |
| fiat | 2.324 |
| seat | 2.070 |
| skoda | 1.631 |
| mazda | 1.549 |
| smart | 1.454 |
| nissan | 1.402 |
| citroen | 1.395 |
| toyota | 1.279 |
| mini | 1.001 |
| volvo | 1.001 |
| hyundai | 0.947 |
| mitsubishi | 0.838 |
| honda | 0.700 |
| sonstige_autos | 0.698 |
| kia | 0.644 |
| suzuki | 0.616 |
| porsche | 0.607 |
| alfa_romeo | 0.572 |
| chevrolet | 0.497 |
| chrysler | 0.352 |
| dacia | 0.287 |
| subaru | 0.261 |
| jeep | 0.212 |
| land_rover | 0.189 |
| jaguar | 0.182 |
| daihatsu | 0.156 |
| saab | 0.152 |
| lancia | 0.131 |
| daewoo | 0.124 |
| rover | 0.124 |
| trabant | 0.100 |
| lada | 0.051 |

```
[47]: print("Total Brands:", cars['brand'].unique().size)
      print("Top 10 brands:")
      cars.groupby(['brand',]).size().sort_values(ascending=False).nlargest(10)
      #The top 10 most sold brands among 40 brands in the market
```

Total Brands: 40

Top 10 brands:

```
[47]: brand
      volkswagen      9162
      bmw             4878
      opel            4499
      mercedes_benz   4142
      audi            3993
      ford            2821
      renault         1945
      peugeot         1324
      fiat            996
      seat            887
      dtype: int64
```

```
[48]: #Variable Model
      print("Total Brands:",cars['model'].unique().size)
      cars.groupby(['brand','model']).size().sort_values(ascending=False).nlargest(10)
      #There are 248 models among them 10 most sold model by sellers are given below
```

Total Brands: 248

```
[48]: brand      model
      volkswagen  golf      3494
      bmw         3er       2489
      volkswagen  polo      1500
      opel        corsa     1393
               astra      1278
      audi        a4        1235
      volkswagen  passat    1207
      mercedes_benz c_klasse 1046
      bmw         5er       1013
      mercedes_benz e_klasse 908
      dtype: int64
```

```
[49]: luxury_cars=cars[(cars['price']>=40000) & (cars['Age'] >= 20)]
      #working range based on observation
```

```
[50]: summary=luxury_cars.groupby(['brand','model']).size().
      →sort_values(ascending=False)
      summary.loc['Grand Total'] = summary.sum()
      lux_cars=summary.to_frame('count')
      lux_cars
      #Luxury cars with age above 20 years and price above 40000 dollars
```

```
[50]:
      brand      model      count
      porsche     911          23
      mercedes_benz sl           3
      jaguar      others          3
      mercedes_benz others        2
```

| | | |
|---------------|-------------|----|
| ford | mustang | 2 |
| alfa_romeo | spider | 2 |
| volkswagen | transporter | 1 |
| renault | others | 1 |
| mercedes_benz | g_klasse | 1 |
| | c_klasse | 1 |
| fiat | others | 1 |
| bmw | m_reihe | 1 |
| Grand Total | | 41 |

```
[51]: porsche_car=cars[(cars['brand']=='porsche') & (cars['price']>=40000) &
      ↪(cars['Age'] >= 20)]
porsche_car.head(5)
#Porsche 911, a luxury cars, is main vehicle which is having high price despite
  ↪its increasing age
```

```
[51]:      seller offerType  price  abtest vehicleType  gearbox  powerPS  \
3286  private      offer  139000    test      coupe    manual      408
4190  private      offer   80000  control     cabrio  automatic      286
5092  private      offer   69800    test      coupe    manual      272
5283  private      offer   52890    test     cabrio    manual      131
8650  private      offer   44300    test      coupe    manual      272
```

| | model | kilometer | fuelType | brand | notRepairedDamage | postalCode | Age |
|------|-------|-----------|----------|---------|-------------------|------------|--------|
| 3286 | 911 | 100000 | petrol | porsche | no | 10629 | 24.000 |
| 4190 | 911 | 80000 | petrol | porsche | no | 50858 | 23.830 |
| 5092 | 911 | 125000 | petrol | porsche | no | 80802 | 23.580 |
| 5283 | 911 | 150000 | petrol | porsche | no | 41812 | 38.580 |
| 8650 | 911 | 150000 | petrol | porsche | no | 9117 | 23.580 |

```
[52]: luxury_cars.groupby(['vehicleType']).size()
#Coupe and Cabrio are the two most sought vehicle among luxury cars
```


```
[52]: vehicleType
bus      1
cabrio   20
coupe    24
limousine 1
suv      1
dtype: int64
```

```
[53]: cars['vehicleType'].value_counts()
```

```
[53]: limousine      11775
small car      9304
station wagon   8100
bus            3600
cabrio         2794
coupe          2262
suv            1815
```

```
others          327
Name: vehicleType, dtype: int64
```

```
[54]: ax=sns.countplot(x='vehicleType',data=cars,order=cars['vehicleType'].
      →value_counts().index)
      for p in ax.patches:
          ax.annotate('{:.2f}'.format(p.get_height()), (p.get_x()+0.15, p.
      →get_height()+1))
      #Limusine is the highest sold vehicle type and least sold are SUV and others
```



```
[55]: # Variable notRepairedDamage
      # yes- car is damaged but not rectified
      # no- car was damaged but has been rectified
      print(cars['notRepairedDamage'].value_counts())
      print("")
      print(pd.
      →crosstab(cars['notRepairedDamage'],columns='count',normalize=True,colnames=['%'])*100)
      # 90% of cars are damage rectified
```

```
no      32574
yes      3994
Name: notRepairedDamage, dtype: int64
```

| % | count |
|-------------------|--------|
| notRepairedDamage | |
| no | 89.078 |
| yes | 10.922 |

```
[56]: print(cars['fuelType'].value_counts())
print("")
print(pd.
      →crosstab(cars['fuelType'],columns='count',normalize=True,colnames=['%'])*100)
# majority of cars are petrol engine followed by diesel engine cars
```

| | |
|---------|-------|
| petrol | 26549 |
| diesel | 12896 |
| lpg | 690 |
| cng | 70 |
| hybrid | 36 |
| electro | 10 |
| other | 6 |

Name: fuelType, dtype: int64

| % | count |
|----------|--------|
| fuelType | |
| cng | 0.174 |
| diesel | 32.034 |
| electro | 0.025 |
| hybrid | 0.089 |
| lpg | 1.714 |
| other | 0.015 |
| petrol | 65.949 |

```
[57]: print(cars['gearbox'].value_counts())
print("")
print(pd.
      →crosstab(cars['gearbox'],columns='count',normalize=True,colnames=['%'])*100)
# 77% of cars have manual gear box
```

| | |
|-----------|-------|
| manual | 32650 |
| automatic | 9410 |

Name: gearbox, dtype: int64

| % | count |
|-----------|--------|
| gearbox | |
| automatic | 22.373 |
| manual | 77.627 |